

Using LDA for Topic Modeling on Song Lyrics

Chenye Yang, Hanchu Zhou, Haodong Liang, Yibo Ma

June 13, 2024

1 Introduction

Analyzing the thematic content of song lyrics presents a fascinating and challenging problem in the field of natural language processing (NLP). Song lyrics can be regarded as a unique form of text that often convey deep emotions and complex narratives within a highly structured and sometimes repetitive linguistic framework. Unlike more straightforward text data, lyrics can feature abstract metaphors, colloquial language, and varying levels of ambiguity, all of which are set to rhythm and melody that influence their interpretation. Such analysis can benefit a range of applications, from enhancing music recommendation systems to aiding in the study of cultural and societal trends through music.

Topic modeling has been widely used in various fields to uncover hidden structures in text data. Blei et al. (2003) introduced latent Dirichlet allocation (LDA) [1], which is a generative probabilistic model that allows sets of text elements to be explained by unobserved groups, providing a way to uncover the hidden thematic structure in a large collection of texts. LDA has since become a standard technique for discovering topics in large corpora. Previous studies have applied LDA to analyze scientific publications, news articles, and social media posts. The potential of LDA to analyze song lyrics, however, remains relatively unexplored.

While LDA is adept at modeling textual data, Dictionary Learning [2, 3] works as a complementary technique, particularly effective in fields requiring robust feature extraction such as image and signal processing. This method involves decomposing a dataset into a dictionary of basis functions and sparse representations, where each data element is expressed as a sparse linear combination of these basis atoms. The focus in Dictionary Learning is on reconstructing the input data with high fidelity, employing a minimal number of active dictionary elements to achieve efficient and compact data representations. The choice of LDA and Dictionary Learning depends on the nature of the data and the analytical goals, whether it is uncovering thematic patterns in text with LDA or achieving sparse and efficient representations in multimedia data with Dictionary Learning.

In addition to LDA, the Term Frequency-Inverse Document Frequency (TF-IDF) [4] model is another crucial technique in the landscape of text analysis. Unlike LDA, which focuses on discovering latent topics, TF-IDF is a statistical measure used to evaluate how important a word is to a document in a collection or corpus. It is often employed as a weighting factor in searches of information retrieval, text mining, and user modeling. The importance increases proportionally to the number of times a word appears in the document but is offset by the frequency of the word in the corpus,

which helps to adjust for the fact that some words appear more frequently in general.

TF-IDF can be particularly effective in distinguishing the relevance of words in large datasets and has been applied extensively in document classification and clustering, spam filtering, and even in sentiment analysis. Its ability to transform the textual data into a numeric form and assess word relevance through both local and global document contexts makes it a valuable complement to LDA in comprehensive text analysis strategies. The combined use of LDA for topic discovery and TF-IDF for importance weighting could potentially provide deeper insights into the thematic elements of song lyrics, enhancing our understanding of linguistic patterns and cultural influences reflected in music.

In this project, we aim to discover the underlying topics in a dataset of song lyrics using TF-IDF and LDA. The goal is to categorize the lyrics into distinct themes and analyze the distribution of these themes across different songs and artists. We preprocess the lyrics, vectorize them using TF-IDF, and fit the LDA model to identify the topics. The results provide insights into the thematic content of the song lyrics and can be used for further analysis or applications in music recommendation systems.

2 Data Collection

The dataset was compiled using the Genius API [5], which provides access to a vast collection of song lyrics. We selected a diverse range of artists and genres to ensure a representative sample. The dataset includes lyrics from well-known artists in Hip Hop, Country, Pop, Rock, Jazz, and Classic. The lyrics were preprocessed to remove non-English words, stopwords, and special characters, ensuring that only valid English words were included in the analysis. The resulting dataset contains a collection of singers, albums, and song lyrics ready for vectorization and topic modeling, shown in Figure 1.

```
1 category,artist,title,lyrics
2 Hip Hop,Kanye West,Mercy,"502 ContributorsTranslationsItalianoFrançaisMercy Lyrics[Intro: Fuzzy Jones]
3 Well, it is a weepin' and a moanin' and a gnashin' of teeth
4 It is a weepin' and a moanin' and a gnashin' of teeth
5 It is a-when it comes to my sound which is the champion sound
```

Figure 1: Example of Lyrics Dataset

3 Method

3.1 Word Vectorization

Word Vectorization is a crucial step in preparing the lyrics for analysis. The following steps were performed:

3.1.1 Lyrics Preprocessing

- i. **Tokenization:** Splitting the lyrics into individual words.

- ii. **Stopwords Removal:** Removing common words that do not contribute much to the meaning (e.g., "the", "is", "and").
- iii. **Custom Stopwords:** Additional stopwords specific to song lyrics (e.g., "chorus", "yeah", "ooh").
- iv. **Filtering Non-English Words:** Ensuring only valid English words are considered.
- v. **Normalization:** Converting all words to lowercase and removing non-alphabetic characters.

3.1.2 TF-IDF Embedding

We use tf-idf vectorization to convert the preprocessed lyrics into numerical vectors. This is a common transformation used in text analysis to represent the importance of a word in a document relative to a collection of documents. The transformation is given by the following steps:

- i. **Term Frequency (TF):** Count the number of times word t appears in document d , denoted as

$$TF(t, d) = f_{t,d}$$

- ii. **Inverse Document Frequency (IDF):** Calculate the inverse document frequency of word t in the entire corpus, denoted as

$$IDF(t) = \log \left(\frac{N}{df_t} \right)$$

where N is the total number of documents and df_t is the number of documents containing word t .

- iii. **TF-IDF:** The TF-IDF value of word t in document d is given by

$$TFIDF(t, d) = TF(t, d) \times IDF(t)$$

In general, the TF-IDF value increases with the number of times a word appears in a document but is offset by the frequency of the word in the entire corpus, making it highlight the distinctive words that are frequently appeared in a document but not common in the entire corpus.

3.2 Latent Dirichlet Allocation

3.2.1 Model Overview

Latent Dirichlet Allocation (LDA) is a two-layer hierarchical Bayesian model that represents documents as random mixtures of latent topics. Each topic is characterized by a distribution over words. The structure of LDA model is shown in Figure 2.




Figure 2: Latent Dirichlet Allocation (LDA) Model

In the above figure, α and β are hyperparameters of Dirichlet Distribution, N is the number of words in a document, K is the number of topics, M is the number of documents, z is the topic assignment for each word, and w is the word.

The LDA model assumes the following generative process for each document d :

1. Draw $\theta_d \sim \text{Dir}(\alpha)$ for each document d
2. Draw $\phi_k \sim \text{Dir}(\beta)$ for each topic k
3. For each word $w_{d,n}$ in document d :
 - (a) Choose a topic $z_{d,n} \sim \text{Multinomial}(\theta_d)$
 - (b) Choose a word $w_{d,n} \sim \text{Multinomial}(\phi_{z_{d,n}})$

3.2.2 Variational Inference

Estimating the parameters of the Latent Dirichlet Allocation (LDA) model involves inferring the posterior distribution of the latent variables given the observed data. Since exact inference is intractable, approximate methods are used. One of the common approaches for parameter estimation in LDA is Variational Inference.

We approximate the posterior distribution $p(\theta, \phi, z|w, \alpha, \beta)$ with a variational distribution $q(\theta, \phi, z|\gamma, \lambda, \psi)$:

$$q(\theta, \phi, z|\gamma, \lambda, \psi) = q(\theta|\gamma)q(\phi|\lambda) \prod_{d=1}^M \prod_{n=1}^{N_d} q(z_{d,n}|\psi_{d,n})$$

where $q(\theta|\gamma)$ is the variational distribution of θ_d , $q(\phi|\lambda)$ is the variational distribution of ϕ_k , $q(z_{d,n}|\psi_{d,n})$ is the variational distribution of $z_{d,n}$. We assume independence between the latent variables so that the joint distribution can be factorized.

We want to minimize the KL divergence:

$$\begin{aligned} \arg \min_{\gamma, \lambda, \psi} KL(q(\theta, \phi, z|\gamma, \lambda, \psi) || p(\theta, \phi, z|w, \alpha, \beta)) &= \mathbb{E}_q[\log q(\theta, \phi, z)] - \mathbb{E}_q[\log p(\theta, \phi, z|w)] \\ &= \mathbb{E}_q[\log q(\theta, \phi, z)] - \mathbb{E}_q[\log p(w, \theta, \phi, z)] + \log p(w) \end{aligned}$$

We know that minimizing the KL divergence is equivalent to maximizing the Evidence Lower Bound (ELBO). Then we can derive ELBO as the objective function for optimization:

$$\begin{aligned}\mathcal{L}(q) &= \mathbb{E}_q[\log p(w, z, \theta, \phi | \alpha, \beta)] - \mathbb{E}_q[\log q(\theta, \phi, z)] \\ &= \mathbb{E}_q[\log p(w | \phi, z)] + \mathbb{E}_q[\log p(z | \theta)] + \mathbb{E}_q[\log p(\theta | \alpha)] + \mathbb{E}_q[\log p(\phi | \beta)] \\ &\quad - \mathbb{E}_q[\log q(\theta)] - \mathbb{E}_q[\log q(\phi)] - \mathbb{E}_q[\log q(z)]\end{aligned}$$

To optimize the ELBO, we iteratively update the variational parameters γ , λ , and ψ until convergence. Note that Dirichlet distribution is the conjugate prior of the multinomial distribution, so we can derive the optimization steps are as follows:

i. Update ψ :

$$\psi_{dnk} \propto \beta_{kw_{dn}} \exp(\mathbb{E}_q[\log \theta_{dk}])$$

ii. Update γ :

$$\gamma_{dk} = \alpha_k + \sum_{n=1}^{N_d} \psi_{dnk}$$

iii. Expectation of Log Topic Proportions:

$$\mathbb{E}_q[\log \theta_{dk}] = \Psi(\gamma_{dk}) - \Psi\left(\sum_{j=1}^K \gamma_{dj}\right)$$

where $\Psi(\cdot)$ is the digamma function.

iv. Update $\lambda_{k,v}$:

$$\lambda_{kv} = \beta_v + \sum_{d=1}^M \sum_{n=1}^{N_d} \psi_{dnk} \mathbb{I}(w_{dn} = v)$$

v. Expectation of Log Word Distributions:

$$\mathbb{E}_q[\log \phi_{dk}] = \Psi(\lambda_{kv}) - \Psi\left(\sum_{j=1}^V \lambda_{kj}\right)$$

The algorithm is summarized in Algorithm 1.

4 Results

We applied the LDA model to the song lyrics dataset and identified five distinct topics. The topics were labeled based on the most frequent words in each topic, providing an intuitive understanding of the themes represented by the lyrics. The top words for each topic were extracted, and the topic distribution across songs was calculated to analyze the prevalence of each theme in the dataset.

Algorithm 1 Variational Inference for LDA

```
1: Initialize  $\gamma, \lambda$  and  $\psi$  randomly.
2: repeat
3:   for each document  $d$  do
4:     for each word  $w_{dn}$  do
5:       Update  $\psi_{dnk}$  for all topics  $k$ .
6:     end for
7:     Update  $\gamma_{dk}$  for all topics  $k$ .
8:   end for
9:   for each topic  $k$  do
10:    Update  $\lambda_{kv}$  for all words  $v$ .
11:   end for
12:   Compute the ELBO and check for convergence.
13: until convergence
```

4.1 Top Words for Each Topic

The top words for each topic are as follows:

- **Topic 0 (Nature and Journey):** song, day, sky, away, night, way, high, new, thing, right.
- **Topic 1 (Love and Emotions):** love, heart, break, thought, head, home, bridge, instrumental, feel, want.
- **Topic 2 (Life and Perceptions):** life, love, look, tell, long, believe, old, time, feel, mind.
- **Topic 3 (Youth and Adolescence):** baby, little, girl, boy, love, night, think, want, right, hear.
- **Topic 4 (World and Ambition):** man, world, need, want, solo, money, instrumental, way, right, bad.

The top words in each topic provide a snapshot of the themes represented by the lyrics in the dataset. We further examine the weights of each top word in the topic to understand the importance of each word. The result is shown in Table 1.

Topic 0	Score	Topic 1	Score	Topic 2	Score	Topic 3	Score	Topic 4	Score
song	36.36	love	32.46	life	21.38	baby	30.95	man	15.03
day	10.65	heart	12.96	love	18.16	little	22.15	world	14.35
sky	9.25	break	11.76	look	17.04	girl	18.40	need	12.55
away	8.87	thought	9.80	tell	16.01	boy	16.17	want	11.83
night	8.78	head	9.73	long	15.57	love	15.25	solo	11.66
way	7.44	home	9.61	believe	15.11	night	15.11	money	10.61
high	7.05	bridge	8.42	old	12.23	think	14.82	instrumental	8.30
new	6.82	instrumental	8.02	time	11.73	want	14.55	way	7.75
thing	6.15	feel	7.92	feel	10.12	right	12.89	right	6.87
right	5.63	want	7.84	mind	8.49	hear	11.72	bad	6.09

Table 1: Top Words in Each Topic with Weights

As Table 1 shows, each topic is characterized by a set of top words that represent the underlying theme. The weights of each word in the topic provide insights into the importance of the word in the topic. For example, the word "love" has a high weight in Topic 1, indicating that this topic is primarily about love and emotions.

The "Nature and Journey" topic is prevalent in songs that describe personal adventures and experiences in the natural world. The "Love and Emotions" topic is common in romantic songs and those expressing deep emotions. The "Life and Perceptions" topic covers songs that reflect on life's deeper meanings and personal's perceptions. The "Youth and Adolescence" topic includes songs about personal connections, youth, and growing up. The "World and Ambition" topic encompasses songs about material wealth, worldly matters, and personal ambitions.

4.2 Topic Distribution

We are interested in the topic distribution of each artist. For each song, we calculate the topic distribution and then aggregate the distribution for each artist. We compared the topic distributions of Taylor Swift, The Beatles, and Justin Bieber. The result is shown in Figure 3. Taylor Swift's lyrics mostly focus on "Life and Perceptions" and "Youth and Adolescence," with a moderate emphasis on "Love and Emotions". The Beatles displays an emphasis on "Life and Perceptions", and "World and Ambition". Justin Bieber's lyrics strongly focus on "Youth and Adolescence".

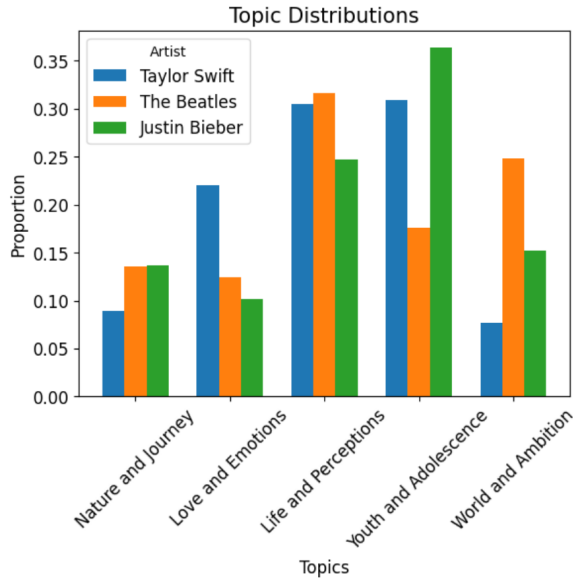


Figure 3: Topic Distribution for Taylor Swift, The Beatles and Justin Bieber

Further, we analyzed the distribution of topics across different music genres. The result is shown in Figure 4. The major topics of Hip Hop are "Life and Perceptions" and "Youth and Adolescence". Country and Pop both have an remarkable emphasis on "Youth and Adolescence". Rock, Jazz and Old Pop, however, have a relatively more balanced representation across all themes, with slightly

different preference over the topics: Rock considers more "World and Ambition", Jazz has more "Nature and Journey" content, and Old Pop sings more "Life and Perceptions".

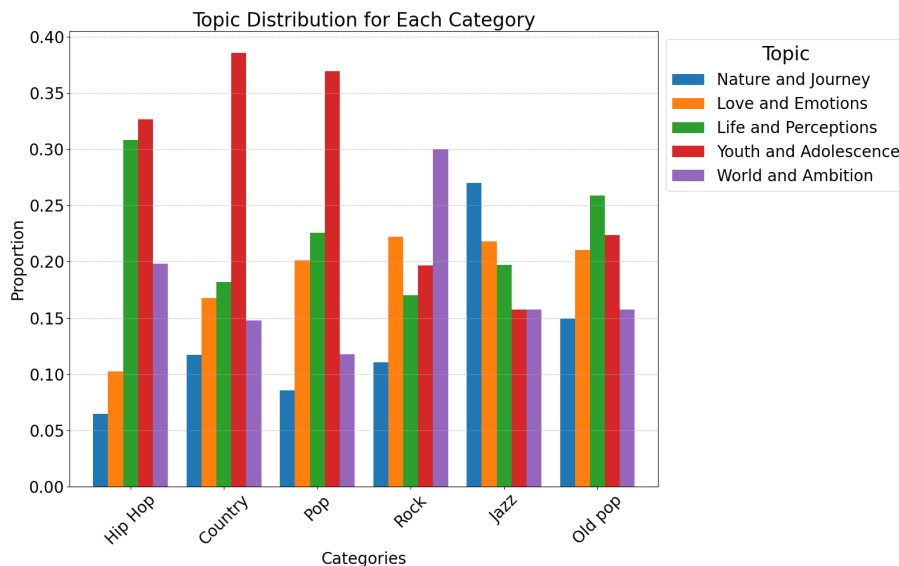


Figure 4: Topic Distribution for Different Music Genres

5 Discussion

The results of the LDA model reveal distinct themes in the song lyrics, reflecting common topics such as love, life, relationships, and experiences. The identification of these themes helps in understanding the emotional and cultural significance of the lyrics. As we mentioned in Section 1, the reason we use LDA, not Dictionary Learning, for lyrics processing is that it excels in extracting thematic content from text, unlike Dictionary Learning, which is better suited for data requiring sparse representation and reconstruction.

6 Conclusion

The LDA model successfully categorized the lyrics into meaningful topics. Each topic represents a distinct theme commonly found in song lyrics, such as love, life, relationships, and experiences. This analysis provides a deeper understanding of the thematic content in song lyrics and can be used for further analysis or applications in music recommendation systems.

7 Future Work

For future work, more advanced models like bidirectional encoder representations from transformers (BERT) [6] can be used to capture deeper contextual relationships in the lyrics, and log-linear

models can be used to reduce computational complexity [7]. Additionally, increasing the number of topics or fine-tuning the preprocessing steps might yield even more precise categorizations. Future studies could also explore the temporal changes in song lyrics topics over the years or across different cultural contexts.

References

- [1] David M Blei, Andrew Y Ng, and Michael I Jordan. “Latent dirichlet allocation”. In: *Journal of machine Learning research* 3.Jan (2003), pp. 993–1022.
- [2] Bruno A Olshausen and David J Field. “Emergence of simple-cell receptive field properties by learning a sparse code for natural images”. In: *Nature* 381.6583 (1996), pp. 607–609.
- [3] Sanjeev Arora et al. “Linear algebraic structure of word senses, with applications to polysemy”. In: *Transactions of the Association for Computational Linguistics* 6 (2018), pp. 483–495.
- [4] Karen Sparck Jones. “A statistical interpretation of term specificity and its application in retrieval”. In: *Journal of documentation* 28.1 (1972), pp. 11–21.
- [5] *Genius API*. URL: <https://genius.com/api-clients>. (accessed: January 2024).
- [6] Jacob Devlin et al. “Bert: Pre-training of deep bidirectional transformers for language understanding”. In: *arXiv preprint arXiv:1810.04805* (2018).
- [7] Tomas Mikolov et al. “Efficient estimation of word representations in vector space”. In: *arXiv preprint arXiv:1301.3781* (2013).

Appendix

```
1 from lyricsgenius import Genius
2 import pandas as pd
3 import numpy as np
4 from langdetect import detect, LangDetectException
5 import nltk
6 from nltk.corpus import stopwords, words
7 from nltk.tokenize import word_tokenize
8 from sklearn.feature_extraction.text import TfidfVectorizer,
   CountVectorizer
9 from sklearn.decomposition import LatentDirichletAllocation
10 import time
11
12 token = "<your-token>"
13 genius = Genius(token)
14
15
16 # Download necessary NLTK data
17 nltk.download('punkt')
18 nltk.download('stopwords')
19
20 # Dictionary of categories and representative singers
21 categories = {
22     "Hip_Hop": ["Kanye_West", "Drake", "Kendrick_Lamar", "Nicki_Minaj",
23                "Cardi_B"],
24     "Country": ["Johnny_Cash", "Dolly_Parton", "Luke_Bryan", "Carrie_
25                Underwood", "Blake_Shelton"],
26     "Pop": ["Taylor_Swift", "Ariana_Grande", "Justin_Bieber", "Katy_
27            Perry", "Lady_Gaga"],
28     "Rock": ["Queen", "Led_Zeppelin", "The_Beatles", "AC/DC", "Pink_
29            Floyd"],
30     "Jazz": ["Louis_Armstrong", "Ella_Fitzgerald", "Duke_Ellington", "
31            Billie_Holiday", "Miles_Davis"],
32     "Old_pop": ["Michael_Jackson", "John_Lennon", "Bob_Marley", "Stevie
33                Wonder", "U2"]
34 }
35
36 # Function to get songs data for an artist
37 def get_songs_data(artist_name, max_songs):
38     artist = genius.search_artist(artist_name, max_songs=max_songs)
39     return [(song.title, song.lyrics) for song in artist.songs]
40
41 # Collect lyrics data
42 all_songs_data = []
```



```

        "whoaaaaaaaa", "uhhhhhhhhhh", "wooooooooo", "wowwwwwwwww", "
        ohhhhhhhhhh", "oohhhhhhhhhh", "laaaaaaaaa", "daaaaaaaaa", "
        naaaaaaaaa", "ehhhhhhhh"]])
62 valid_words = set(words.words())
63
64 def preprocess_lyrics(lyrics):
65     tokens = word_tokenize(lyrics.lower())
66     filtered_tokens = [word for word in tokens if word.isalnum() and
        word not in stop_words and word in valid_words]
67     english_lyrics = []
68     for word in filtered_tokens:
69         if word.isalpha(): # Check if the word is an English
        alphabetic word
70             english_lyrics.append(word)
71     return ' '.join(english_lyrics)
72
73 df['processed_lyrics'] = df['lyrics'].apply(preprocess_lyrics)
74
75 # Vectorize the lyrics
76 vectorizer = TfidfVectorizer(max_features=1000, max_df=0.5, min_df=0.1,
        stop_words='english')
77 dtm = vectorizer.fit_transform(df['processed_lyrics'])
78
79 # Perform LDA
80 lda = LatentDirichletAllocation(n_components=5, random_state=1)
81 lda.fit(dtm)
82
83 # Display the topics
84 def display_topics(model, feature_names, no_top_words):
85     for topic_idx, topic in enumerate(model.components_):
86         print("Topic_%d:" % (topic_idx))
87         print(" ".join([feature_names[i] for i in topic.argsort()[:-
        no_top_words - 1:-1]]))
88
89 no_top_words = 10
90 tf_feature_names = vectorizer.get_feature_names_out()
91 display_topics(lda, tf_feature_names, no_top_words)
92
93 # Function to display the top words and their scores for each topic
94 def display_topics_with_scores(model, feature_names, no_top_words):
95     for topic_idx, topic in enumerate(model.components_):
96         print(f"Topic_{topic_idx}:")
97         top_words_scores = [(feature_names[i], topic[i]) for i in topic
        .argsort()[:-no_top_words - 1:-1]]
98         for word, score in top_words_scores:
99             print(f"{word}: {score:.4f}")

```

```

100         print("\n")
101
102     # Number of top words to display
103     no_top_words = 5
104
105     # Assuming your LDA model and vectorizer are already fitted and named 'lda' and 'vectorizer', respectively
106     # Get feature names (words) from the vectorizer
107     feature_names = vectorizer.get_feature_names_out()
108
109     # Display the topics and their scores
110     display_topics_with_scores(lda, feature_names, no_top_words)
111
112     # Get the topic distribution for each song
113     topic_values = lda.transform(dtm)
114
115     # Rename the topics based on the summaries
116     topic_names = {
117         0: "Journey/Experience",
118         1: "Love/Desire",
119         2: "Life/Contemplation",
120         3: "Relationships/Belief",
121         4: "Wealth/World"
122     }
123
124     # Create a DataFrame to display the topic values for each song
125     topic_values_df = pd.DataFrame(topic_values, columns=[topic_names[i]
126         for i in range(lda.n_components)])
127     result_df = pd.concat([df[['artist', 'title']], topic_values_df], axis
128         =1)
129
130     # Display the topic values for all songs
131     print(result_df.head())
132
133     # Save the topic values to a CSV file
134     result_df.to_csv("topic_distribution.csv", index=False)
135
136     import matplotlib.pyplot as plt
137     topic_names = ["Nature_and_Journey", "Love_and_Emotions", "Life_and_Perceptions", "Youth_and_Adolescence", "World_and_Ambition"]
138
139     def get_topic_distribution(artist_name):
140         # Filter the artist's songs
141         artist_songs = df[df['artist'] == artist_name]
142         artist_songs['processed_lyrics'] = artist_songs['lyrics'].apply(
143             preprocess_lyrics)

```

```

141
142     # Vectorize the artist's lyrics using the same vectorizer
143     artist_dtm = vectorizer.transform(artist_songs['processed_lyrics'])
144
145     # Use the trained LDA model to get topic distributions
146     artist_topic_distributions = lda.transform(artist_dtm)
147
148     # Convert topic distributions to a DataFrame
149     artist_topic_vectors_df = pd.DataFrame(artist_topic_distributions,
150                                             columns=[f"Topic_{i}" for i in range(num_topics)])
151
152     # Calculate the mean topic distribution
153     artist_mean_topic_distribution = artist_topic_vectors_df.mean()
154
155     return artist_mean_topic_distribution
156
157 # Get topic distributions for Taylor Swift and The Beatles
158 taylor_topic_distribution = get_topic_distribution('Taylor_Swift')
159 beatles_topic_distribution = get_topic_distribution('The_Beatles')
160 justin_topic_distribution = get_topic_distribution('Justin_Bieber')
161
162 # Combine the topic distributions into one DataFrame
163 combined_df = pd.DataFrame({
164     'Taylor_Swift': taylor_topic_distribution,
165     'The_Beatles': beatles_topic_distribution,
166     'Justin_Bieber': justin_topic_distribution
167 })
168
169 # Plot the combined topic distributions
170 plt.figure(figsize=(16, 20))
171 combined_df.plot(kind='bar', width=0.7)
172 plt.title("Topic_Distributions", fontsize=15)
173 plt.xlabel("Topics", fontsize=12)
174 plt.ylabel("Proportion", fontsize=12)
175 plt.xticks(ticks=range(num_topics), labels=topic_names, rotation=45,
176           fontsize=12)
177 plt.yticks(fontsize=12)
178 plt.legend(title="Artist", fontsize=12)
179 plt.show()
180
181 import matplotlib.pyplot as plt
182
183 # Define summarized topic names
184 topic_names = ["Nature_and_Journey", "Love_and_Emotions", "Life_and_
185               Perceptions", "Youth_and_Adolescence", "World_and_Ambition"]

```

```

184 def get_category_topic_distribution(category_name):
185     # Filter the category's songs
186     category_songs = df[df['category'] == category_name]
187     category_songs['processed_lyrics'] = category_songs['lyrics'].apply
        (preprocess_lyrics)
188
189     # Vectorize the category's lyrics using the same vectorizer
190     category_dtm = vectorizer.transform(category_songs['
        processed_lyrics'])
191
192     # Use the trained LDA model to get topic distributions
193     category_topic_distributions = lda.transform(category_dtm)
194
195     # Convert topic distributions to a DataFrame
196     category_topic_vectors_df = pd.DataFrame(
        category_topic_distributions, columns=[f"Topic_{i}" for i in
        range(num_topics)])
197
198     # Calculate the mean topic distribution
199     category_mean_topic_distribution = category_topic_vectors_df.mean()
200
201     return category_mean_topic_distribution
202
203 # List of categories
204 categories = df['category'].unique()
205
206 # Get topic distributions for each category
207 category_distributions = {category: get_category_topic_distribution(
        category) for category in categories}
208
209 # Combine the topic distributions into one DataFrame
210 combined_category_df = pd.DataFrame(category_distributions).T
211 combined_category_df.columns = topic_names
212
213 # Plot the combined topic distributions
214 plt.figure(figsize=(14, 8))
215 combined_category_df.plot(kind='bar', width=0.8, figsize=(16, 10))
216 plt.title("Topic_Distribution_for_Each_Category", fontsize=24)
217 plt.xlabel("Categories", fontsize=20)
218 plt.ylabel("Proportion", fontsize=20)
219 plt.xticks(rotation=45, fontsize=20)
220 plt.yticks(fontsize=20)
221 plt.legend(title="Topic", fontsize=20, title_fontsize=24, loc='upper_
        left', bbox_to_anchor=(1, 1))
222 plt.grid(axis='y', linestyle='—', alpha=0.7)
223 plt.tight_layout()

```

```
224 plt.savefig('genres.png')
225 plt.show()
```