

## EEC 289A Assignment 2 Report

Chenye Yang, Hanchu Zhou, Haodong Liang, Yibo Ma

### 1 Introduction

Texture synthesis has been an active research topic in computer vision both as a way to verify texture analysis methods, as well as in its own right. Potential applications of a successful texture synthesis algorithm are broad, including occlusion fill-in, lossy image and video compression, foreground removal, etc.

This project seeks to reproduce and expand upon the work by Efros & Leung (1999) on non-parametric texture synthesis [1]. Originally, this method was groundbreaking for its ability to synthesize new texture fields that maintain the visual characteristics of a given sample texture. Our primary objective is to apply this technique to the textures assigned in the class, exploring both the effectiveness and the limitations of the method in replicating and scaling up these specific patterns.

In addition to replication, this project will extend the methodology to other images and contexts of interest. Particularly, we aim to test the feasibility of applying texture synthesis techniques to generate MNIST digits and a galaxy picture<sup>1</sup>, which present a unique challenge in scaling the approach to handle more complex and structured data.

### 2 Algorithm Overview

This algorithm synthesizes texture by expanding pixel by pixel from an initial seed point. It uses a single pixel,  $p$ , as the synthesis unit to capture as much high-frequency information as possible. All previously synthesized pixels within a square window surrounding  $p$  (with weights to highlight local structures) serve as the context. To continue the synthesis, probability tables for the distribution of  $p$  are needed for all possible contexts. However, while these tables are typically manageable in size for textual data, constructing them explicitly for textures is impractical. An approximation might be achieved using various clustering techniques, but we opt not to construct a model at all. Instead, each new context queries the sample image, and the distribution of  $p$  is created as a histogram of all possible values found in the sample image, as demonstrated in Figure 1. This non-parametric sampling method, though simple, is highly effective at capturing statistical processes for which an adequate model has not yet been developed.

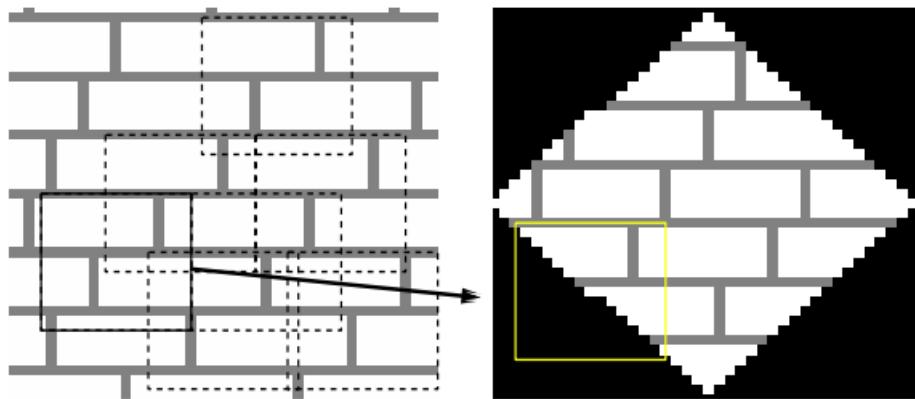


Figure 1: Algorithm Overview

In this work, the texture is modeled as a Markov Random Field (MRF). That is, we assume that the probability distribution of brightness values for a pixel, given the brightness values of its spatial neighborhood, is independent of the rest of the image. The neighborhood of a pixel is modeled as a square window around

<sup>1</sup>Galaxy starry night sky background, free public domain CC0 photo. <https://www.rawpixel.com/image/5924106/photo-image-public-domain-stars-galaxy>

that pixel. The size of the window is a free parameter that specifies how stochastic the user believes this texture to be. More specifically, if the texture is presumed to be mainly regular at high spatial frequencies and mainly stochastic at low spatial frequencies, the size of the window should be on the scale of the biggest regular feature.

## 2.1 Synthesizing one pixel

Let  $I$  be an image that is being synthesized from a texture sample image  $I_{\text{smp}} \subset I_{\text{real}}$  where  $I_{\text{real}}$  is the real infinite texture. Let  $p \in I$  be a pixel and let  $\omega(p) \subset I$  be a square image patch of width  $w$  centered at  $p$ . Let  $d_{\text{perc}}(\omega_1, \omega_2)$  denote some perceptual distance between two patches. Assume that all pixels in  $I$  except for  $p$  are known. To synthesize the value of  $p$  we first construct an approximation to the conditional probability distribution  $P(p | \omega(p))$  and then sample from it.

Based on our MRF model, we assume that  $p$  is independent of  $I \setminus \omega(p)$  given  $\omega(p)$ . Define a set

$$\Omega(p) = \{\omega' \subset I_{\text{real}} : d_{\text{perc}}(\omega_0, \omega(p)) = 0\}$$

containing all occurrences of  $\omega(p)$  in  $I_{\text{real}}$ , then the conditional PDF of  $p$  can be estimated with a histogram of all center pixel values in  $\Omega(p)$ . Unfortunately, we are only given  $I_{\text{smp}}$ , a finite sample from  $I_{\text{real}}$ , which means there might not be any matches for  $\omega(p)$  in  $I_{\text{smp}}$ . Thus, we must use a heuristic which will let us find a plausible  $\Omega'(p) \approx \Omega(p)$  to sample from. In our implementation, a variation of the k-Nearest Neighbors technique is used: the closest match  $\omega'_{\text{best}} = \arg \min_{\omega} d_{\text{perc}}(\omega(p), \omega) \in I_{\text{smp}}$  is found, and all image patches  $\omega_0$  with  $d_{\text{perc}}(\omega'_{\text{best}}, \omega') < \epsilon$  are included in  $\Omega'$ , where  $\epsilon$  is a threshold. The center pixel values of patches in  $\Omega'$  give us a histogram for  $p$ , which can then be sampled, either uniformly or weighted by  $d_{\text{perc}}$ .

Now it only remains to find a suitable  $d_{\text{perc}}$ . One choice is a normalized sum of squared differences metric  $d_{\text{SSD}}$ . However, this metric gives the same weight to any mismatched pixel, whether near the center or at the edge of the window. Since we would like to preserve the local structure of the texture as much as possible, the error for nearby pixels should be greater than for pixels far away. To achieve this effect we set  $d_{\text{perc}} = d_{\text{SSD}} \cdot G$  where  $G$  is a two-dimensional Gaussian kernel.

## 2.2 Synthesizing texture

In the previous section, we discussed a method of synthesizing a pixel when its neighborhood pixels are already known. Unfortunately, this method cannot be used for synthesizing the entire texture or even for hole-filling (unless the hole is just one pixel) since for any pixel the values of only some of its neighborhood pixels will be known. The correct solution would be to consider the joint probability of all pixels together, but this is intractable for images of realistic size.

Instead, a Shannon-inspired heuristic is proposed, where the texture is grown in layers outward from a 3-by-3 seed taken randomly from the sample image (in case of hole-filling, the synthesis proceeds from the edges of the hole). Now for any point  $p$  to be synthesized only some of the pixel values in  $\omega(p)$  are known (i.e., have already been synthesized). Thus, the pixel synthesis algorithm must be modified to handle unknown neighborhood pixel values. This can be easily done by only matching on the known values in  $\omega(p)$  and normalizing the error by the total number of known pixels when computing the conditional pdf for  $p$ . This heuristic does not guarantee that the pdf for  $p$  will stay valid as the rest of  $\omega(p)$  is filled in. However, it appears to be a good approximation in practice. One can also treat this as an initialization step for an iterative approach such as Gibbs sampling. However, our trials have shown that Gibbs sampling produced very little improvement for most textures. This lack of improvement indicates that the heuristic indeed provides a good approximation to the desired conditional pdf.

# 3 Experiments

## 3.1 Textures

We apply the non-parametric texture synthesis method to the following textures: the textures provided in the class, the MNIST digits, and the galaxy picture we chosen. The textures provided in the class have the size around  $530 \times 530$ , shown in Figure 2.



(a) Texture 2



(b) Texture 4



(c) Texture 9

Figure 2: Texture: provided in the class

For the MNIST digits, we load the MNIST dataset ( $28 \times 28$  handwritten digits)<sup>2</sup>, including 60000 training images and 10000 testing images. We randomly select  $20 \times 20$  images from the training set as the sample texture of size  $560 \times 560$ , shown in Figure 3.

```
1 5 2 7 9 6 6 5 5 9 1 1 4 6 9 3 4 9 0
7 8 7 5 7 8 7 7 0 4 5 1 2 9 4 4 1 8 9 8
0 3 1 3 1 5 2 2 2 9 8 7 5 5 8 3 0 1 6 2
3 9 7 4 0 4 7 7 8 3 1 0 0 3 3 0 7 2 2 7
0 6 7 0 8 0 1 9 0 2 9 9 0 7 2 0 1 2 9 5
0 2 1 6 9 8 2 9 3 4 1 1 0 2 9 2 2 0 6 6
6 2 3 3 4 0 7 2 0 1 2 4 6 2 1 5 8 9 8 5
6 8 3 7 7 9 4 1 3 9 6 7 0 5 0 6 0 9 7 9
1 6 1 5 1 1 3 5 0 6 9 7 5 1 8 5 6 8 3 1
9 8 6 5 5 3 2 4 1 2 3 6 1 3 5 0 1 1 9 1
4 2 4 1 2 5 9 1 2 2 5 4 1 7 1 5 6 5 9 6
3 2 3 6 1 3 6 4 8 6 8 8 5 8 1 4 5 2 3 8
2 1 6 3 0 0 2 9 0 1 4 0 1 1 6 4 7 2 5 2
2 4 6 0 5 5 9 4 8 1 6 4 0 2 2 6 8 9 2 0
4 9 4 8 0 4 2 9 8 6 4 1 3 6 6 1 0 7 3 5
9 7 7 0 0 2 1 2 6 8 1 8 8 6 1 8 4 3 8 1
5 7 8 5 5 7 9 0 7 4 2 0 2 6 6 5 0 9 9 6
5 8 1 2 1 1 5 6 3 0 2 5 3 8 3 5 3 6 8 0
6 4 5 5 2 7 1 1 5 1 0 2 5 1 2 9 7 7 1 2
9 7 2 6 8 6 4 0 3 7 3 2 0 6 0 4 2 4 7 5
```

Figure 3: Texture: MNIST digits

For the galaxy picture, we load the galaxy picture ( $800 \times 585$ ), shown in Figure 4.



Figure 4: Texture: galaxy picture

<sup>2</sup>Downloaded from [https://git-disl.github.io/GTDLBench/datasets/mnist\\_datasets/](https://git-disl.github.io/GTDLBench/datasets/mnist_datasets/)

### 3.2 Non-parametric Texture Synthesis

The algorithm proposed by Efros & Leung (1999) [1], Algorithm 1, synthesizes texture by modeling the image as a Markov Random Field (MRF) where each pixel's value is dependent only on the values in its immediate neighborhood. The texture synthesis process grows the new image pixel by pixel from an initial seed based on this model. Here are some important features about the algorithm:

- Non-parametric approach: The algorithm does not build an explicit model but estimates pixel distributions dynamically from the sample image.
- Control over randomness: The size of the neighborhood window can be adjusted to control the randomness of the texture, affecting how structured or stochastic the synthesized texture appears.
- Efficiency: A variation of the k Nearest Neighbors technique is employed to find matching neighborhoods quickly.
- Local structure preservation: Emphasis on maintaining the local structure of the texture to ensure visual continuity and realism in the synthesized image.

---

**Algorithm 1** Texture Synthesis by Non-parametric Sampling

---

- 1: Initialize with a small seed taken randomly from the sample image  $I_{\text{smp}}$ .
  - 2: **while** there are unsynthesized pixels in the image  $I$  **do**
  - 3:     Select an unsynthesized pixel  $p$  in  $I$ .
  - 4:     Define the neighborhood  $N(p)$  of  $p$  as a square window centered at  $p$ .
  - 5:     Find all neighborhoods in  $I_{\text{smp}}$  that are similar to  $N(p)$  using a perceptual distance metric  $d_{\text{perc}}$ .
  - 6:     Construct a histogram of pixel values from these similar neighborhoods to approximate the conditional probability distribution  $P(p|N(p))$ .
  - 7:     Sample from this distribution to set the value of  $p$ .
  - 8: **end while**
- 

### 3.3 Parallelization: Multi-core Processing

To speed up the texture synthesis process, we parallelize the algorithm by dividing the image into blocks and synthesizing each block independently, given that the processing of each pixel (and its surrounding window) is somewhat independent. We use the Python `multiprocessing` library to create a pool of worker processes that can work on different blocks concurrently:

- Split the work: Distribute the processing of each pixel to different CPU cores. This can be done by creating a list of tasks where each task contains the information needed to process a pixel.
- Process pool: Use a multiprocessing pool to process these tasks concurrently.
- Result collection: Collect the results from each task and update the main data structures accordingly.

With this parallelization strategy, we can achieve a significant speedup (depending on hardware, around  $8\times$ ) in the texture synthesis process.

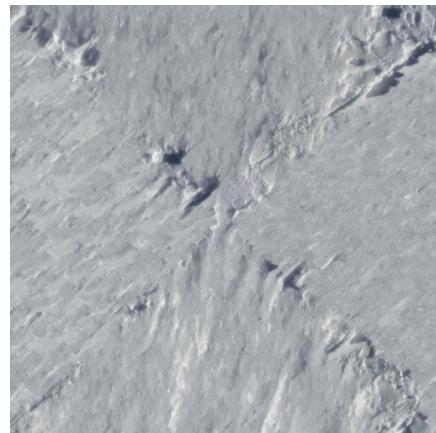
## 4 Results

### 4.1 Texture Synthesis

In this section, we present the results of applying the non-parametric texture synthesis method to the textures described in the previous section. The patch size is set to  $11 \times 11$  for all the textures, and the resolution is set to  $1000 \times 1000$  for all synthesized pictures. The synthesized textures are shown in Figure 5 - 9.



(a) Texture 2

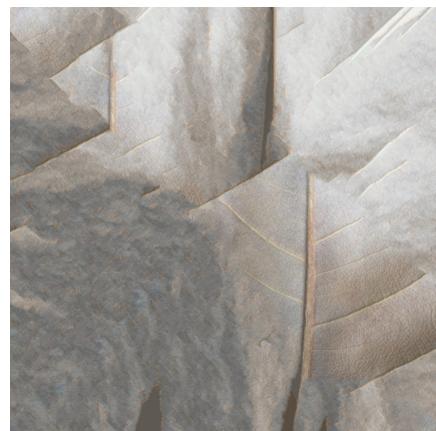


(b) Synthesized Texture 2

Figure 5: Synthesis of Texture 2



(a) Texture 4



(b) Synthesized Texture 4

Figure 6: Synthesis of Texture 4



(a) Texture 9



(b) Synthesized Texture 9

Figure 7: Synthesis of Texture 9

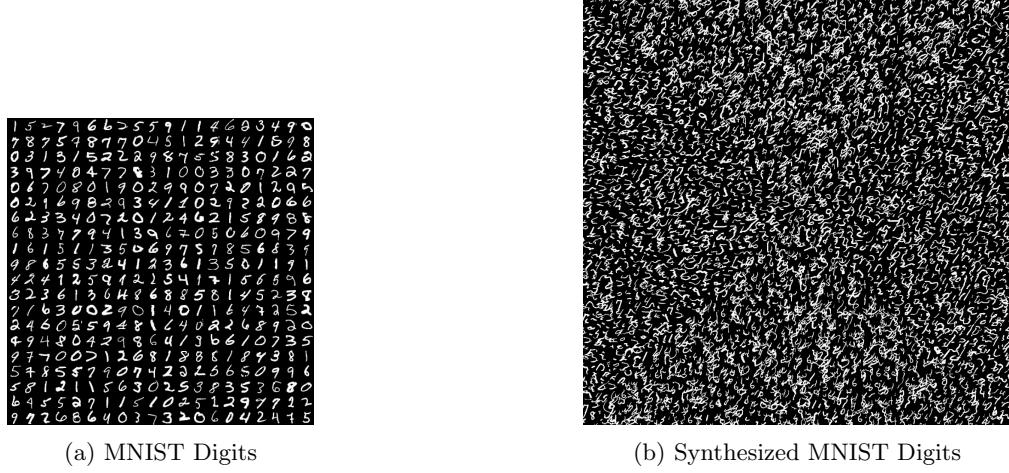


Figure 8: Synthesis of MNIST Digits

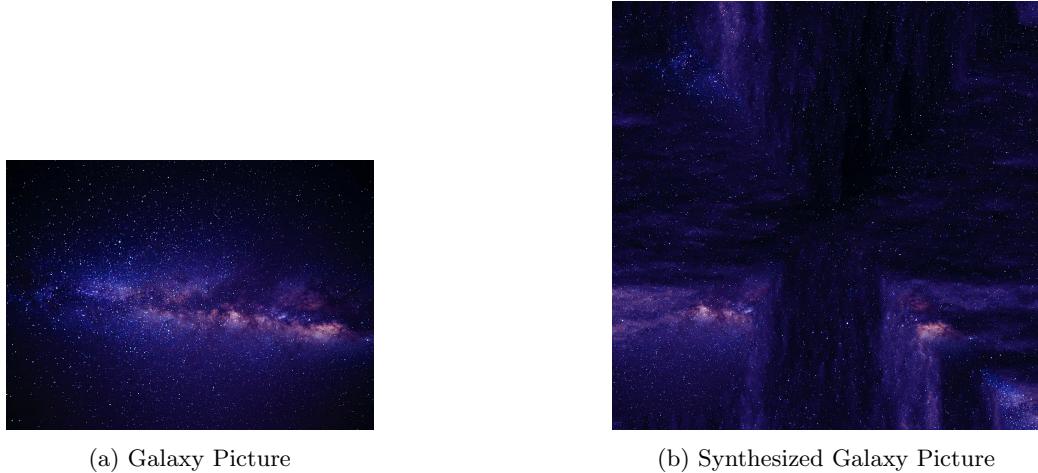


Figure 9: Synthesis of Galaxy Picture

## 4.2 Different Neighborhood Patch Sizes

In this section, we present the results of applying the non-parametric texture synthesis method on MNIST with different neighborhood sizes / patch sizes. The results are shown in Figure 10.

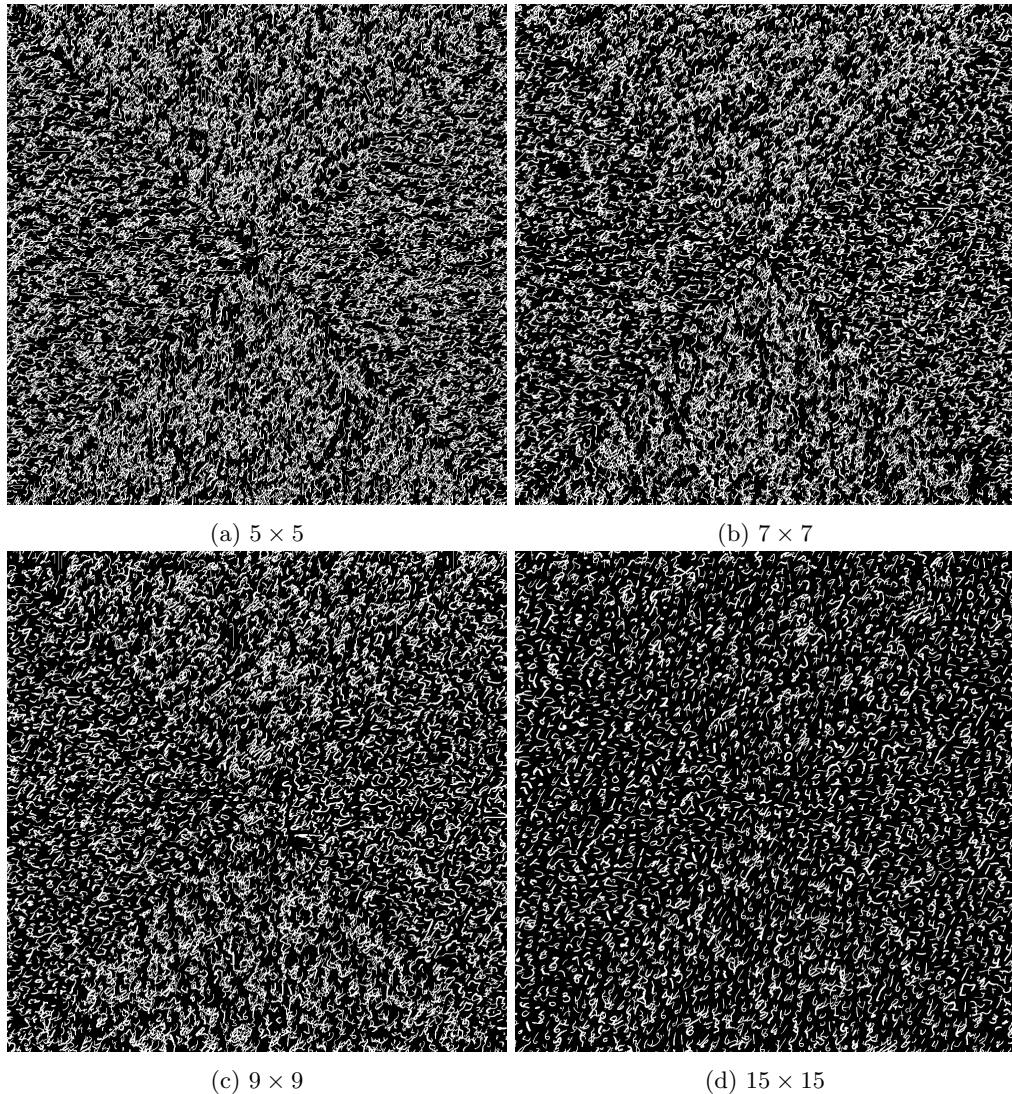


Figure 10: Different Neighborhood Sizes for MNIST

As we can observe from Figure 10, when we scale the method to synthesize MNIST digits, we need to adjust the patch size to appropriately utilize the neighborhood information. In the above four cases, the patch size of  $15 \times 15$  produces the best synthesized result.

## References

- [1] A.A. Efros and T.K. Leung. "Texture synthesis by non-parametric sampling". In: *Proceedings of the Seventh IEEE International Conference on Computer Vision*. Vol. 2. 1999, 1033–1038 vol.2. DOI: 10.1109/ICCV.1999.790383.

## Appendix

We referred to the GitHub repo: <https://github.com/goldbema/TextureSynthesis>, modified the code with `multiprocessing` parallelization, and tried PyTorch matrix operation speedup.

```
1  , ,
2  This module is a Python implementation of:
3
4  A. Efros and T. Leung, "Texture Synthesis by Non-parametric Sampling,"  

5  Proceedings of the Seventh IEEE International Conference on Computer  

6  Vision, September 1999.
7
8  Specifically, this module implements texture synthesis by growing a 3x3  

9  texture patch  

10 pixel-by-pixel. Please see the authors' project page for additional algorithm  

11 details:  

12
13 Example:  

14
15 Generate a 50x50 texture patch from a texture available at the input path  

16 and save it to  

17 the output path. Also, visualize the synthesis process:  

18
19
20
21
22 --author-- = 'Maxwell_Goldberg'  

23 --author-- = 'Chenyue_Yang, _Hanchu_Zhou, _Haodong_Liang, _Yibo_Ma'  

24
25 import torch
26 import time
27 import torch.nn.functional as F
28 import argparse
29 import cv2
30 import numpy as np
31 from multiprocessing import Pool
32
33 EIGHT_CONNECTED_NEIGHBOR KERNEL = np.array([[1., 1., 1.],  

34                                         [1., 0., 1.],  

35                                         [1., 1., 1.]], dtype=np.float64)
36 SIGMA_COEFF = 6.4      # The denominator for a 2D Gaussian sigma used in the  

37                         # reference implementation.
38 ERROR_THRESHOLD = 0.1  # The default error threshold for synthesis acceptance  

39                         # in the reference implementation.
40
41 # def normalized_ssd(sample, window, mask):
42 #     if torch.backends.mps.is_available():
43 #         device = torch.device("mps")
```

```
44 #     elif torch.cuda.is_available():
45 #         device = torch.device("cuda")
46 #     else:
47 #         device = torch.device("cpu")
48
49 #     # Convert inputs to PyTorch tensors and move them to the specified
50 #     # device
51 #     sample = torch.tensor(sample, dtype=torch.float32).to(device)
52 #     window = torch.tensor(window, dtype=torch.float32).to(device)
53 #     mask = torch.tensor(mask, dtype=torch.float32).to(device)
54
55 #     # Get dimensions
56 #     sh, sw = sample.shape
57 #     wh, ww = window.shape
58
59 #     # Compute 2D Gaussian kernel
60 #     sigma = wh / SIGMA_COEFF # SIGMA_COEFF should be defined elsewhere in
61 #     # your code
62 #     kernel = torch.Tensor(cv2.getGaussianKernel(ksize=wh, sigma=sigma)).to(
63 #     device)
64 #     kernel_2d = torch.mm(kernel, kernel.t())
65
66 #     # Apply the Gaussian kernel to the mask
67 #     weighted_mask = mask * kernel_2d
68
69 #     # Calculate padded size for valid convolution
70 #     padded_sample = F.pad(sample, (ww//2, ww//2, wh//2, wh//2))
71
72 #     # Perform convolution using the flipped window (correlation)
73 #     window = window.flip([0, 1])
74 #     ssd_map = F.conv2d(padded_sample[None, None, :, :], window[None, None,
75 #     :, :], None, stride=1)
76
77 #     # Compute sum of squared differences
78 #     squared_diff = (ssd_map - torch.sum(window)**2)**2
79
80 #     # Multiply by the weighted mask and sum over the kernel
81 #     result_map = F.conv2d(squared_diff, weighted_mask[None, None, :, :],
82 #     None, stride=1)
83
84 #     # Normalize the SSD by the maximum possible contribution
85 #     total_ssd = torch.sum(weighted_mask)
86 #     normalized_ssd_map = result_map / total_ssd
87
88 #     return normalized_ssd_map[0, 0].cpu().numpy()
89
90
91
92 # def normalized_ssd(sample, window, mask):
93 #     if torch.backends.mps.is_available():
94 #         device = torch.device("mps")
95 #     elif torch.cuda.is_available():
96 #         device = torch.device("cuda")
97 #     else:
98 #         device = torch.device("cpu")
```

```

93  #     # Ensure all inputs are float32
94  #     sample, window, mask = map(lambda x: torch.tensor(x, dtype=torch.float32
95  # , device=device), (sample, window, mask))
96  #
97  #     # Calculate Gaussian kernel in float32, ensure vectors are 1D by
98  #     # flattening
99  #     sigma = window.shape[0] / SIGMA_COEFF
100 #     k1 = torch.tensor(cv2.getGaussianKernel(ksize=window.shape[0], sigma=
101 #     sigma).flatten(), dtype=torch.float32, device=device)
102 #     kernel = torch.outer(k1, k1)
103 #
104 #     # Use unfold to create sliding windows
105 #     unfolded_sample = F.unfold(sample[None, None], kernel_size=window.shape,
106 #     padding=0, stride=1)
107 #     unfolded_sample = unfolded_sample.view(1, window.numel(), sample.shape
108 #     [0]-window.shape[0]+1, sample.shape[1]-window.shape[1]+1)
109 #
110 #     ssd = (unfolded_sample - window.view(1, -1, 1, 1))**2
111 #     ssd *= kernel.view(1, -1, 1, 1) * mask.view(1, -1, 1, 1)
112 #     ssd = ssd.sum(dim=1)
113 #
114 #     # Normalize SSD
115 #     total_ssd = torch.sum(mask * kernel)
116 #     normalized_ssd = ssd / total_ssd
117 #     return normalized_ssd.cpu().numpy()
118
119 def normalized_ssd(sample, window, mask):
120     wh, ww = window.shape
121     sh, sw = sample.shape
122
123     # Get sliding window views of the sample, window, and mask.
124     strided_sample = np.lib.stride_tricks.as_strided(sample, shape=((sh-wh+1),
125         (sw-ww+1), wh, ww),
126             strides=(sample.strides[0], sample.strides[1], sample.
127                 strides[0], sample.strides[1]))
128     strided_sample = strided_sample.reshape(-1, wh, ww)
129
130     # Note that the window and mask views have the same shape as the strided
131     # sample, but the kernel is fixed
132     # rather than sliding for each of these components.
133     strided_window = np.lib.stride_tricks.as_strided(window, shape=((sh-wh+1),
134         (sw-ww+1), wh, ww),
135             strides=(0, 0, window.strides[0], window.strides[1]))
136     strided_window = strided_window.reshape(-1, wh, ww)
137
138     strided_mask = np.lib.stride_tricks.as_strided(mask, shape=((sh-wh+1),
139         (sw-ww+1), wh, ww),
140             strides=(0, 0, mask.strides[0], mask.strides[1]))
141     strided_mask = strided_mask.reshape(-1, wh, ww)
142
143     # Form a 2D Gaussian weight matrix from symmetric linearly separable
144     # Gaussian kernels and generate a
145     # strided view over this matrix.
146     sigma = wh / SIGMA_COEFF

```

```

136     kernel = cv2.getGaussianKernel(ksize=wh, sigma=sigma)
137     kernel_2d = kernel * kernel.T
138
139     strided_kernel = np.lib.stride_tricks.as_strided(kernel_2d, shape=((sh-wh
140         +1), (sw-ww+1), wh, ww),
141                     strides=(0, 0, kernel_2d.strides[0], kernel_2d.strides
142                     [1]))
143     strided_kernel = strided_kernel.reshape(-1, wh, ww)
144
145     # Take the sum of squared differences over all sliding sample windows and
146     # weight it so that only existing neighbors
147     # contribute to error. Use the Gaussian kernel to weight central values
148     # more strongly than distant neighbors.
149     squared_differences = ((strided_sample - strided_window)**2) *
150         strided_kernel * strided_mask
151     ssd = np.sum(squared_differences, axis=(1,2))
152     ssd = ssd.reshape(sh-wh+1, sw-ww+1)
153
154     # Normalize the SSD by the maximum possible contribution.
155     total_ssd = np.sum(mask * kernel_2d)
156     normalized_ssd = ssd / total_ssd
157
158     return normalized_ssd
159
160
161 def get_candidate_indices(normalized_ssd, error_threshold=ERROR_THRESHOLD):
162     min_ssd = np.min(normalized_ssd)
163     min_threshold = min_ssd * (1. + error_threshold)
164     indices = np.where(normalized_ssd <= min_threshold)
165     return indices
166
167 def select_pixel_index(normalized_ssd, indices, method='uniform'):
168     N = indices[0].shape[0]
169
170     if method == 'uniform':
171         weights = np.ones(N) / float(N)
172     else:
173         weights = normalized_ssd[indices]
174         weights = weights / np.sum(weights)
175
176     # Select a random pixel index from the index list.
177     selection = np.random.choice(np.arange(N), size=1, p=weights)
178     selected_index = (indices[0][selection], indices[1][selection])
179
180     return selected_index
181
182 def get_neighboring_pixel_indices(pixel_mask):
183     # Taking the difference between the dilated mask and the initial mask
184     # gives only the 8-connected neighbors of the mask frontier.
185     kernel = np.ones((3,3))
186     dilated_mask = cv2.dilate(pixel_mask, kernel, iterations=1)
187     neighbors = dilated_mask - pixel_mask
188
189     # Recover the indices of the mask frontier.

```

```

185     neighbor_indices = np.nonzero(neighbors)
186
187     return neighbor_indices
188
189 def permute_neighbors(pixel_mask, neighbors):
190     N = neighbors[0].shape[0]
191
192     # Generate a permutation of the neighboring indices
193     permuted_indices = np.random.permutation(np.arange(N))
194     permuted_neighbors = (neighbors[0][permuted_indices], neighbors[1][
195         permuted_indices])
196
196     # Use convolution to count the number of existing neighbors for all
197     # entries in the mask.
197     neighbor_count = cv2.filter2D(pixel_mask, ddepth=-1, kernel=
198         EIGHT_CONNECTED_NEIGHBOR_KERNEL, borderType=cv2.BORDER_CONSTANT)
199
200     # Sort the permuted neighboring indices by quantity of existing neighbors
201     # descending.
202     permuted_neighbor_counts = neighbor_count[permuted_neighbors]
203
204     sorted_order = np.argsort(permuted_neighbor_counts)[::-1]
205     permuted_neighbors = (permuted_neighbors[0][sorted_order],
206                           permuted_neighbors[1][sorted_order])
207
207     return permuted_neighbors
208
209
210 def texture_can_be_synthesized(mask):
211     # The texture can be synthesized while the mask has unfilled entries.
212     mh, mw = mask.shape[:2]
213     num_completed = np.count_nonzero(mask)
214     num_incomplete = (mh * mw) - num_completed
215
216     return num_incomplete > 0
217
218
219 def initialize_texture_synthesis(original_sample, window_size, kernel_size):
220     # Convert original to sample representation.
221     sample = cv2.cvtColor(original_sample, cv2.COLOR_BGR2GRAY)
222
223     # Convert sample to floating point and normalize to the range [0., 1.]
224     sample = sample.astype(np.float64)
225     sample = sample / 255.
226
227     # Generate window
228     window = np.zeros(window_size, dtype=np.float64)
229
230     # Generate output window
231     if original_sample.ndim == 2:
232         result_window = np.zeros_like(window, dtype=np.uint8)
233     else:
234         result_window = np.zeros(window_size + (3,), dtype=np.uint8)
235
235     # Generate window mask
236     h, w = window.shape

```

```

234     mask = np.zeros((h, w), dtype=np.float64)
235
236     # Initialize window with random seed from sample
237     sh, sw = original_sample.shape[:2]
238     ih = np.random.randint(sh-3+1)
239     iw = np.random.randint(sw-3+1)
240     seed = sample[ih:ih+3, iw:iw+3]
241
242     # Place seed in center of window
243     ph, pw = (h//2)-1, (w//2)-1
244     window[ph:ph+3, pw:pw+3] = seed
245     mask[ph:ph+3, pw:pw+3] = 1
246     result_window[ph:ph+3, pw:pw+3] = original_sample[ih:ih+3, iw:iw+3]
247
248     # Obtain padded versions of window and mask
249     win = kernel_size//2
250     padded_window = cv2.copyMakeBorder(window,
251                                         top=win, bottom=win, left=win, right=
252                                         win, borderType=cv2.BORDER_CONSTANT
253                                         , value=0.)
254     padded_mask = cv2.copyMakeBorder(mask,
255                                         top=win, bottom=win, left=win, right=win,
256                                         borderType=cv2.BORDER_CONSTANT,
257                                         value=0.)
258
259     # Obtain views of the padded window and mask
260     window = padded_window[win:-win, win:-win]
261     mask = padded_mask[win:-win, win:-win]
262
263     return sample, window, mask, padded_window, padded_mask, result_window
264
265 def process_pixel(args):
266     sample, window_slice, mask_slice, kernel_size, ch, cw, original_sample =
267         args
268     ssd = normalized_ssd(sample, window_slice, mask_slice)
269     indices = get_candidate_indices(ssd)
270     selected_index = select_pixel_index(ssd, indices)
271     selected_index = (selected_index[0] + kernel_size // 2, selected_index[1]
272                     + kernel_size // 2)
273     return ch, cw, sample[selected_index], original_sample[selected_index[0],
274                         selected_index[1]]
275
276 def synthesize_texture(original_sample, window_size, kernel_size, visualize):
277     global gif_count
278     (sample, window, mask, padded_window, padded_mask, result_window) =
279         initialize_texture_synthesis(original_sample, window_size, kernel_size
280 )
281
282     pool = Pool()
283
284     while texture_can_be_synthesized(mask):
285         neighboring_indices = get_neighboring_pixel_indices(mask)
286         neighboring_indices = permute_neighbors(mask, neighboring_indices)

```

```
279     tasks = []
280     for ch, cw in zip(neighboring_indices[0], neighboring_indices[1]):
281         window_slice = padded_window[ch:ch+kernel_size, cw:cw+kernel_size]
282         mask_slice = padded_mask[ch:ch+kernel_size, cw:cw+kernel_size]
283         tasks.append((sample, window_slice, mask_slice, kernel_size, ch,
284                       cw, original_sample))
285
286     results = pool.map(process_pixel, tasks)
287
288     for ch, cw, new_value, result_value in results:
289         window[ch, cw] = new_value
290         mask[ch, cw] = 1
291         result_window[ch, cw] = result_value
292
293         if visualize:
294             cv2.imshow('synthesis_window', result_window)
295             key = cv2.waitKey(1)
296             if key == 27:
297                 cv2.destroyAllWindows()
298                 pool.close()
299                 pool.join()
300                 return result_window
301
302     pool.close()
303     pool.join()
304
305     if visualize:
306         cv2.imshow('synthesis_window', result_window)
307         cv2.waitKey(0)
308         cv2.destroyAllWindows()
309
310     return result_window
311
312 def validate_args(args):
313     wh, ww = args.window_height, args.window_width
314     if wh < 3 or ww < 3:
315         raise ValueError('window_size must be greater than or equal to (3,3).')
316
317     if args.kernel_size <= 1:
318         raise ValueError('kernel_size must be greater than 1.')
319
320     if args.kernel_size % 2 == 0:
321         raise ValueError('kernel_size must be odd.')
322
323     if args.kernel_size > min(wh, ww):
324         raise ValueError('kernel_size must be less than or equal to the smaller window_size dimension.')
325
326 def parse_args():
327     parser = argparse.ArgumentParser(description='Perform texture synthesis')
328     parser.add_argument('--sample_path', type=str, required=True, help='Path to the texture sample')
```

```
329     parser.add_argument('--out_path', type=str, required=False, help='Output_'
330                         _path_for_synthetized_texture')
331     parser.add_argument('--window_height', type=int, required=False, default=
332                         50, help='Height_of_the_synthetization_window')
333     parser.add_argument('--window_width', type=int, required=False, default=
334                         50, help='Width_of_the_synthetization_window')
335     parser.add_argument('--kernel_size', type=int, required=False, default=11,
336                         help='One_dimension_of_the_square_synthetization_kernel')
337     parser.add_argument('--visualize', required=False, action='store_true',
338                         help='Visualize_the_synthetization_process')
339     args = parser.parse_args()
340     return args
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357 def main():
    args = parse_args()

    sample = cv2.imread(args.sample_path)
    if sample is None:
        raise ValueError('Unable_to_read_image_from_sample_path.')
    validate_args(args)

    start_time = time.time()
    synthesized_texture = synthesize_texture(original_sample=sample,
                                                window_size=(args.window_height,
                                                              args.window_width),
                                                kernel_size=args.kernel_size,
                                                visualize=args.visualize)
    print('Synthesis_time: {:.2f} seconds'.format(time.time() - start_time))

    if args.out_path is not None:
        cv2.imwrite(args.out_path, synthesized_texture)

if __name__ == '__main__':
    main()
```